

Homework 1

April 17, 2019

1 Homework 1

1.1 Initial Set-Up:

```
In [1]: # Read in data manipulation packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Use for data-preprocessing
wine_dat = pd.read_csv('./wine_modified.csv')
print('Wine modified dimensions:', wine_dat.shape)
wine_dat.head()

# Use for remaining sections
X_train = pd.read_csv('./wine_train_data.csv')
y_train = pd.read_csv('./wine_train_labels.csv')
X_valid = pd.read_csv('./wine_val_data.csv')
y_valid = pd.read_csv('./wine_val_labels.csv')
X_test = pd.read_csv('./wine_test_data.csv')
y_test = pd.read_csv('./wine_test_labels.csv')
print('Traning data dimensions:', X_train.shape)
print('Number of rows of training:', X_train.shape[0])
print('Number of rows of validation:', X_valid.shape[0])
print('Number of rows of testing:', X_test.shape[0])
```

Wine modified dimensions: (178, 14)

Traning data dimensions: (100, 13)

Number of rows of training: 100

Number of rows of validation: 39

Number of rows of testing: 39

1.2 Data Preprocessing:

Problem 1:

```
In [4]: # Where are the missing values?
wine_dat.isnull().sum()
```

```

Out[4]: class          10
        Alcohol        0
        Malic acid     12
        Ash            112
        Alcalinity of ash 12
        Magnesium      22
        Total phenols   0
        Flavanoids     48
        Nonflavanoid phenols 14
        Proanthocyanins 12
        Color intensity 12
        Hue            14
        OD280/OD315     11
        Proline         14
        dtype: int64

```

```

In [25]: # Remove NaN in Class
        wine_clean = wine_dat.dropna(subset=['class'])
        # Remove Rows with > 7 missing
        wine_clean = wine_clean.dropna(thresh=wine_clean.shape[1]-7)
        # Report new number of rows
        print('New number of rows:', wine_clean.shape[0])

```

New number of rows: 154

Problem 2:

```

In [26]: # Detect breakdown of missing values
        wine_clean.isnull().sum()/wine_clean.shape[0]

```

```

Out[26]: class          0.000000
        Alcohol        0.000000
        Malic acid     0.000000
        Ash            0.616883
        Alcalinity of ash 0.000000
        Magnesium      0.058442
        Total phenols   0.000000
        Flavanoids     0.227273
        Nonflavanoid phenols 0.000000
        Proanthocyanins 0.000000
        Color intensity 0.000000
        Hue            0.000000
        OD280/OD315     0.000000
        Proline         0.000000
        dtype: float64

```

```

In [27]: # Remove ash
        wine_clean.drop('Ash', axis=1, inplace=True)

```

```

# Print new shape
print('New data dimensions:', wine_clean.shape)
print('Removed Ash feature')

```

New data dimensions: (154, 13)

Removed Ash feature

```

In [28]: # Fill in missing data with mean
features_detect = wine_clean.columns[pd.isnull(wine_clean).sum() > 0].tolist()
print(features_detect, 'are features with missing values')
wine_clean = wine_clean.fillna(wine_clean.mean())
print('New standard deviations are:')
wine_clean[features_detect].std()

```

['Magnesium', 'Flavanoids'] are features with missing values

New standard deviations are:

```

Out[28]: Magnesium      14.440377
         Flavanoids      0.873573
         dtype: float64

```

Problem 3:

```

In [31]: # Remove rows that are within 3 standard deviations (~99.7%) of the data.
wine_clean = wine_clean[np.abs(wine_clean.Alcohol - wine_clean.Alcohol.mean()) <= (3*wine_clean.Alcohol.std())]
wine_clean = wine_clean[np.abs(wine_clean.Proline - wine_clean.Proline.mean()) <= (3*wine_clean.Proline.std())]
print('Final dimensions are:', wine_clean.shape)

```

Final dimensions are: (148, 13)

While there are many different ways to classify an individual data point as an outlier, the most straightforward and indisputable approach is to remove rows where the value is above or below 3 standard deviations of the data. This value would need to be more extreme than 99.7% of the rest of the data. Hence, the data point is pretty safe to be classified as an outlier and removed.

1.3 Decision Trees:

Problem 4:

```

In [4]: from sklearn.tree import DecisionTreeClassifier, export_graphviz
import pydotplus
from sklearn.neighbors import KNeighborsClassifier
from sklearn import preprocessing

# For final predictions:
X_train_fin = pd.concat([X_train, X_valid], ignore_index=True)

```

```

y_train_fin = pd.concat([y_train,y_valid],ignore_index=True)

best_acc = 0
computations = ['gini', 'entropy']

for i in computations:
    clf = DecisionTreeClassifier(i)
    clf.fit(X_train, y_train)
    accuracy = np.sum(clf.predict(X_valid)==y_valid['class'].values)*1/len(y_valid['class'].values)
    print('For criterion =', i, 'the validation accuracy = ' + str(accuracy))
    if (accuracy > best_acc):
        best_acc = accuracy
        best_computation = i
print("")
print('The best criterion is', best_computation, 'with a validation accuracy of', best_acc)

clf = DecisionTreeClassifier(criterion=best_computation)
clf.fit(X_train_fin, y_train_fin)
predictions = clf.predict(X_test)
fin_accuracy = np.sum(predictions==y_test['class'].values)*1/len(y_test['class'].values)
print('\n')
print('Final accuracy on total training data = ' + str(fin_accuracy))

```

For criterion = gini the validation accuracy = 0.948717948718

For criterion = entropy the validation accuracy = 0.974358974359

The best criterion is entropy with a validation accuracy of 0.974358974359

Final accuracy on total training data = 0.820512820513

Problem 5:

```

In [13]: best_acc = 0
         min_sample = [2,5,10,20]

         for i in min_sample:
             clf = DecisionTreeClassifier(best_computation,min_samples_split=i)
             clf.fit(X_train, y_train)
             accuracy = np.sum(clf.predict(X_valid)==y_valid['class'].values)*1/len(y_valid['class'].values)
             print('For min sample =', i, 'the validation accuracy = ' + str(accuracy))
             if (accuracy > best_acc):
                 best_acc = accuracy
                 best_sample = i
         print("")
         print('The best sample is', best_sample, 'with a validation accuracy of', best_acc)

```

```

clf = DecisionTreeClassifier(criterion=best_computation,min_samples_split = best_sample)
clf.fit(X_train_fin, y_train_fin)
predictions = clf.predict(X_test)
fin_accuracy = np.sum(predictions==y_test['class'].values)*1/len(y_test['class'].values)
print('\n')
print('Final accuracy = ' + str(fin_accuracy))

```

For min sample = 2 the validation accuracy = 0.948717948718
 For min sample = 5 the validation accuracy = 0.948717948718
 For min sample = 10 the validation accuracy = 0.923076923077
 For min sample = 20 the validation accuracy = 0.948717948718

The best sample is 2 with a validation accuracy of 0.948717948718

Final accuracy = 0.820512820513

Problem 6:

```

In [15]: sample_size = [20,40,60,80,100]
        acc_list = []

        for i in sample_size:
            clf = DecisionTreeClassifier(best_computation,min_samples_split=best_sample)
            clf.fit(X_train[:i], y_train[:i])
            accuracy = np.sum(clf.predict(X_valid)==y_valid['class'].values)*1/len(y_valid['class'].values)
            print('For training data sample size =', i, 'the validation accuracy = ' + str(accuracy))
            acc_list.append(accuracy)

```

For training data sample size = 20 the validation accuracy = 0.641025641026
 For training data sample size = 40 the validation accuracy = 0.846153846154
 For training data sample size = 60 the validation accuracy = 0.871794871795
 For training data sample size = 80 the validation accuracy = 0.871794871795
 For training data sample size = 100 the validation accuracy = 0.948717948718

```


In [21]: # Generate Line Plot
        plt.plot(sample_size,acc_list)

        # Label the axes
        plt.xlabel('Sample Size')
        plt.ylabel('Accuracy')

        #Figure Title
        plt.title('More Training Data Achieves Greater Accuracy')

        plt.show()

```



1.4 KNN

Problem 7:

```
In [22]: # Use Standard Scaler from Sklearn to normalize data in full training data:
normalize = preprocessing.StandardScaler().fit(X_train_fin)
X_train_fin_norm = normalize.transform(X_train_fin)

# Apply normalizing of aggregate training set to other data-sets:
X_train_norm = normalize.transform(X_train)
X_valid_norm = normalize.transform(X_valid)
X_test_norm = normalize.transform(X_test)

In [28]: # Build KNN classifier
clf = KNeighborsClassifier(n_neighbors=3, metric = 'euclidean')
clf.fit(X_train_fin_norm, y_train_fin.values.ravel())
predictions = clf.predict(X_test_norm)
accuracy = np.sum(predictions == y_test['class'].values)*1.0/len(y_test['class'].values)
print('Test accuracy with k = 3 is', accuracy)
```

Test accuracy with k = 3 is 0.871794871795

Problem 8:

```

In [31]: best_acc = 0
         dist_metric = ['euclidean', 'manhattan', 'chebyshev']

         for i in dist_metric:
             clf = KNeighborsClassifier(n_neighbors=3, metric = i)
             clf.fit(X_train_norm, y_train)
             accuracy = np.sum(clf.predict(X_valid_norm)==y_valid['class'].values)*1/len(y_valid)
             print('For distance metric =', i, 'the validation accuracy = ' + str(accuracy))
             if (accuracy > best_acc):
                 best_acc = accuracy
                 best_dist = i

         print("")
         print('The best distance metric is', best_dist, 'with a validation accuracy of', best_a

         clf = KNeighborsClassifier(n_neighbors=3, metric = best_dist)
         clf.fit(X_train_fin_norm, y_train_fin.values.ravel())
         predictions = clf.predict(X_test_norm)
         fin_accuracy = np.sum(predictions==y_test['class'].values)*1/len(y_test['class'].values)
         print('\n')
         print('Final accuracy = ' + str(fin_accuracy))

```

For distance metric = euclidean the validation accuracy = 0.923076923077

For distance metric = manhattan the validation accuracy = 0.948717948718

For distance metric = chebyshev the validation accuracy = 0.923076923077

The best distance metric is manhattan with a validation accuracy of 0.948717948718

Final accuracy = 0.974358974359

/Users/kkannapp/anaconda/lib/python3.6/site-packages/ipykernel_launcher.py:6: DataConversionWarn

Problem 9:

```

In [32]: k_list = [1,3,5,7,9]
         acc_list = []
         best_acc = 0

         for i in k_list:
             clf = KNeighborsClassifier(n_neighbors=i, metric = 'euclidean')
             clf.fit(X_train_norm, y_train)
             accuracy = np.sum(clf.predict(X_valid_norm)==y_valid['class'].values)*1/len(y_valid)
             print('For k =', i, 'the validation accuracy = ' + str(accuracy))
             acc_list.append(accuracy)
             if (accuracy > best_acc):
                 best_acc = accuracy

```

```

        best_k = i
    print("")
    print('The best k-value is', best_k, 'with a validation accuracy of', best_acc)

    clf = KNeighborsClassifier(n_neighbors=best_k, metric = 'euclidean')
    clf.fit(X_train_fin_norm, y_train_fin.values.ravel())
    predictions = clf.predict(X_test_norm)
    fin_accuracy = np.sum(predictions==y_test['class'].values)*1/len(y_test['class'].values)
    print('\n')
    print('Final accuracy = ' + str(fin_accuracy))

```

```

For k = 1 the validation accuracy = 0.948717948718
For k = 3 the validation accuracy = 0.923076923077
For k = 5 the validation accuracy = 0.948717948718
For k = 7 the validation accuracy = 0.974358974359
For k = 9 the validation accuracy = 0.948717948718

```

The best k-value is 7 with a validation accuracy of 0.974358974359

Final accuracy = 0.923076923077

```

/Users/kkannapp/anaconda/lib/python3.6/site-packages/ipykernel_launcher.py:7: DataConversionWarn
import sys

```

```


In [34]: # Generate Line Plot
         plt.plot(k_list, acc_list)

         # Label the axes
         plt.xlabel('Value of K')
         plt.ylabel('Accuracy')

         #Figure Title
         plt.title('Value of K has Large Impact on Performance')

         plt.show()

```

Problem 10:

```
In [36]: sample_size = [20,40,60,80,100]
        acc_list = []

        for i in sample_size:
            clf = KNeighborsClassifier(n_neighbors=3,metric = 'euclidean')
            clf.fit(X_train_norm[:i], y_train[:i])
            accuracy = np.sum(clf.predict(X_valid_norm)==y_valid['class'].values)*1/len(y_valid)
            print('For sample size =', i, 'the validation accuracy = ' + str(accuracy))
            acc_list.append(accuracy)
```

For sample size = 20 the validation accuracy = 0.948717948718

For sample size = 40 the validation accuracy = 1.0

For sample size = 60 the validation accuracy = 1.0

For sample size = 80 the validation accuracy = 1.0

For sample size = 100 the validation accuracy = 0.923076923077

/Users/kkannapp/anaconda/lib/python3.6/site-packages/ipykernel_launcher.py:6: DataConversionWarn

```
In [37]: # Generate Line Plot
        plt.plot(sample_size,acc_list)
```

```
# Label the axes
plt.xlabel('Sample Size')
plt.ylabel('Accuracy')

#Figure Title
plt.title('Shuffling Data Important, Predictions could be Unrealistic')

plt.show()
```

