

Homework 4

February 25, 2019

0.1 Worksheet 8 - Generative models 3

Problem 3:

Part a)

Unpacked and sorted through the directories, have 20 classifications of news types, which are informed by the directory hierarchy

Part b)

Leverage data-set and hierarchy on scikit-learn. Links to the same directory as specified in the homework. A training label will link fo the lableed value according to the directory that it lies in, i.e.: alt.atheism 1, etc.

```
In [11]: from sklearn.datasets import fetch_20newsgroups
         from sklearn.feature_extraction.text import TfidfVectorizer
         # Remove strong identifiers of article category
         newsgroups_train = fetch_20newsgroups(subset='train',remove=('headers', 'footers', 'quotes'))
         # Remove strong identifiers of article category
         newsgroups_test = fetch_20newsgroups(subset='test',remove=('headers', 'footers', 'quotes'))

In [2]: print(newsgroups_train.filenames.shape)
        print(newsgroups_test.filenames.shape)

(11314,)
(7532,)
```

We have 11,314 documents of training data and 7,532 documents of test data.

Part c)

```
In [3]: length = newsgroups_train.filenames.shape[0]

In [4]: import numpy as np
         unique, counts = np.unique(newsgroups_train.target, return_counts=True)

In [5]: import pandas as pd
         prior_prob = pd.DataFrame({'class':unique, 'prior_prob':counts/length})

In [6]: # Fraction of total documents that belong to each class. Appear to be less on the last class
         # Appears that the class is transformed from the range 0-19 as opposed to 1-20.
         prior_prob
```

```
Out [6]:
```

	class	prior_prob
0	0	0.042425
1	1	0.051617
2	2	0.052236
3	3	0.052148
4	4	0.051087
5	5	0.052413
6	6	0.051706
7	7	0.052501
8	8	0.052855
9	9	0.052766
10	10	0.053032
11	11	0.052590
12	12	0.052236
13	13	0.052501
14	14	0.052413
15	15	0.052943
16	16	0.048259
17	17	0.049850
18	18	0.041100
19	19	0.033322

```
In [8]: vocab = {}
reverse_vocab = {}
count = 0
a = open('./vocabulary.txt', 'r')
for v in a:
    val = v.strip()
    vocab[val] = count
    reverse_vocab[count] = val
    count += 1
```

```
In [252]: vocab['baseball']
```

```
Out [252]: 1816
```

```
In [215]: # Vectorize each training document, using the vocabulary document
vectorizer = TfidfVectorizer(strip_accents='unicode', decode_error = 'ignore', stop_wo
vectors = vectorizer.fit_transform(newsgroups_train.data)
vectors.shape
```

```
Out [215]: (11314, 61188)
```

```
In [216]: # Vectorize the test document, using the vocabulary document
vectors_test = vectorizer.fit_transform(newsgroups_test.data)
vectors_test.shape
```

```
Out [216]: (7532, 61188)
```

Part d) Used a different smoothing constant than 1, 1 did not perform as well. MultinomialNB uses logs inherently.

```
In [223]: from sklearn.naive_bayes import MultinomialNB
          from sklearn import metrics
          clf = MultinomialNB(alpha=0.046)
          clf.fit(vectors, newsgroups_train.target)
          # Naive bayes uses prior probability distributions from above
```

```
Out[223]: MultinomialNB(alpha=0.046, class_prior=None, fit_prior=True)
```

Part e)

```
In [224]: # Predictions:
          pred = clf.predict(vectors_test)
```

```
In [225]: print('The model is', round(metrics.accuracy_score(newsgroups_test.target, pred)*100),
          print('The model has an error rate of', round((1- metrics.accuracy_score(newsgroups_test.target, pred)*100), 1))
```

The model is 70.0 % accurate

The model has an error rate of 30.0 %

0.2 Worksheet 9 - Clustering

Problem 1:

```
In [6]: f = open('./Animals_with_Attributes/Features/README-features.txt', 'r')
        file_contents = f.read()
        print (file_contents)
        f.close()
```

```
=====
Pre-Extracted Image Features for the Animals_with_Attributes Dataset
=====
```

Names and IDs of all classes are in the file

- classes.txt

Predicate names and IDs are in the file

- predicates.txt

Training and test classnames for the attribute-based classifier are in

- trainclasses.txt, testclasses.txt

The class->attribute matrix is given in three formats.

- predicate-matrix-numeric.txt (positive numeric entries, some missing entries encoded as -1)

- predicate-matrix-binary.txt (binarized with mean of feature, missing

set to 0)
- predicate-matrix.png (PNG image file for visual inspection)

Note that the entries are in the order of the names/predicates files,
not alphabetically.

There are 6 feature representations:

- cq: (global) color histogram (1x1 + 2x2 + 4x4 spatial pyramid, 128 bins each, each histogram L1-normalized)
- lss[1]: local self similarity (2000 entry codebook, raw bag-of-visual-word counts)
- phog[2]: histogram of oriented gradients (1x1 + 2x2 + 4x4 spatial pyramid, 12 bins each, each histogram L1-normalized or all zero)
- rgsift[3]: rgSIFT descriptors (2000 entry codebook, bag-of-visual-word counts, L1-normalized)
- sift[4]: SIFT descriptors (2000 entry codebook, raw bag-of-visual-word counts)
- surf[5]: SUFT descriptors (2000 entry codebook, raw bag-of-visual-word counts)

Each file consists of one sample per line in ASCII format.
All representations have non-negative entries.

Instructions for fixed-split Attribute-Based Classification

- train using all examples of all 40 classes from "trainclasses.txt"
- test on all examples of 10 classes specified in the "testclasses.txt"

Instructions for CV-like Multi-Class Classification

Follow a protocol similar to Caltech256:

- choose N_{train} \in {5,10,15,20,25,30,40,50} and set $N_{\text{test}}=25$
- train on N_{train} random training examples from each class
- test on N_{test} random examples out of the remaining images from each class
- calculate the mean of the confusion matrix (class averaged accuracy)
- report averaged results for 10-times this procedure

References

- [1] E. Shechtman, and M. Irani: "Matching Local Self-Similarities across Images and Videos", CVPR 2007.

- [2] A. Bosch, A. Zisserman, and X. Munoz: "Representing shape with a spatial pyramid kernel", CIVR 2007.
- [3] Koen E. A. van de Sande, Theo Gevers and Cees G. M. Snoek: "Evaluation of Color Descriptors for Object and Scene Recognition", CVPR 2008.
- [4] D. G. Lowe, "Distinctive Image Features from Scale-Invariant Keypoints", IJCV 2004.
- [5] H. Bay, T. Tuytelaars, and L. Van Gool: "SURF: Speeded Up Robust Features", ECCV 2006.

Problem 2:

```
In [7]: # Different animal classes
        f = open('./Animals_with_Attributes/classes.txt', 'r')
        classes_str = f.read()
        print (classes_str)
        f.close()

1      antelope
2      grizzly+bear
3      killer+whale
4      beaver
5      dalmatian
6      persian+cat
7      horse
8      german+shepherd
9      blue+whale
10     siamese+cat
11     skunk
12     mole
13     tiger
14     hippopotamus
15     leopard
16     moose
17     spider+monkey
18     humpback+whale
19     elephant
20     gorilla
21     ox
22     fox
23     sheep
24     seal
25     chimpanzee
```

```
26      hamster
27      squirrel
28      rhinoceros
29      rabbit
30      bat
31      giraffe
32      wolf
33      chihuahua
34      rat
35      weasel
36      otter
37      buffalo
38      zebra
39      giant+panda
40      deer
41      bobcat
42      pig
43      lion
44      mouse
45      polar+bear
46      collie
47      walrus
48      raccoon
49      cow
50      dolphin
```

```
In [8]: # Different available features
        f = open('./Animals_with_Attributes/predicates.txt', 'r')
        features_str = f.read()
        print (features_str)
        f.close()
```

```
1      black
2      white
3      blue
4      brown
5      gray
6      orange
7      red
8      yellow
9      patches
10     spots
11     stripes
12     furry
13     hairless
14     toughskin
```

15	big
16	small
17	bulbous
18	lean
19	flippers
20	hands
21	hooves
22	pads
23	paws
24	longleg
25	longneck
26	tail
27	chewteeth
28	meatteeth
29	buckteeth
30	strainteeth
31	horns
32	claws
33	tusks
34	smelly
35	flys
36	hops
37	swims
38	tunnels
39	walks
40	fast
41	slow
42	strong
43	weak
44	muscle
45	bipedal
46	quadrapedal
47	active
48	inactive
49	nocturnal
50	hibernate
51	agility
52	fish
53	meat
54	plankton
55	vegetation
56	insects
57	forager
58	grazer
59	hunter
60	scavenger
61	skimmer
62	stalker

```
63     newworld
64     oldworld
65     arctic
66     coastal
67     desert
68     bush
69     plains
70     forest
71     fields
72     jungle
73     mountains
74     ocean
75     ground
76     water
77     tree
78     cave
79     fierce
80     timid
81     smart
82     group
83     solitary
84     nestspot
85     domestic
```

Problem 3:

```
In [9]: classes = ''.join([i for i in classes_str if not i.isdigit()]).split()
```

```
In [10]: classes[:10]
```

```
Out[10]: ['antelope',
          'grizzly+bear',
          'killer+whale',
          'beaver',
          'dalmatian',
          'persian+cat',
          'horse',
          'german+shepherd',
          'blue+whale',
          'siamese+cat']
```

```
In [11]: features = ''.join([i for i in features_str if not i.isdigit()]).split()
```

```
In [12]: features[:10]
```

```
Out[12]: ['black',
          'white',
```



```

'blue',
'brown',
'gray',
'orange',
'red',
'yellow',
'patches',
'spots']

```

```

In [13]: animals_data = pd.read_fwf("./Animals_with_Attributes/predicate-matrix-continuous.txt",
    print('The shape of the data is', animals_data.shape)

```

The shape of the data is (50, 85)

```

In [14]: animal_df = pd.DataFrame(data = animals_data, columns = features)
    animal_df.index = classes
    animal_df.head()

```

```

Out[14]:

```

	black	white	blue	brown	gray	orange	red	yellow	patches	\
antelope	-1.00	-1.00	-1.0	-1.00	12.34	0.0	0.0	0.0	16.11	
grizzly+bear	39.25	1.39	0.0	74.14	3.75	0.0	0.0	0.0	1.25	
killer+whale	83.40	64.79	0.0	0.00	1.25	0.0	0.0	0.0	68.49	
beaver	19.38	0.00	0.0	87.81	7.50	0.0	0.0	0.0	0.00	
dalmatian	69.58	73.33	0.0	6.39	0.00	0.0	0.0	0.0	37.08	

	spots	...	water	tree	cave	fierce	timid	smart	\
antelope	9.19	...	0.00	0.00	1.23	10.49	39.24	17.57	
grizzly+bear	0.00	...	7.64	9.79	53.14	61.80	12.50	24.00	
killer+whale	32.69	...	79.49	0.00	0.00	38.27	9.77	52.03	
beaver	7.50	...	65.62	0.00	0.00	3.75	31.88	41.88	
dalmatian	100.00	...	1.25	6.25	0.00	9.38	31.67	53.26	

	group	solitary	nestspot	domestic
antelope	50.59	2.35	9.70	8.38
grizzly+bear	3.12	58.64	20.14	11.39
killer+whale	24.94	15.77	13.41	15.42
beaver	23.44	31.88	33.44	13.12
dalmatian	24.44	29.38	11.25	72.71

[5 rows x 85 columns]

```

In [15]: # Import K Means Package
    from sklearn.cluster import KMeans

    # Set k = 10
    km10 = KMeans(n_clusters=10)
    km10.fit(animals_data)
    # Get cluster assignment labels

```

```

labels = km10.labels_
# Format results as a DataFrame
results = pd.DataFrame([animal_df.index, labels]).T
results.columns = ['class', 'cluster']

```

```
In [16]: results.groupby('cluster')['class'].apply(list)
```

```

Out[16]: cluster
0          [spider+monkey, gorilla, chimpanzee]
1          [hippopotamus, elephant, rhinoceros]
2          [grizzly+bear, polar+bear]
3          [antelope, horse, giraffe, zebra, deer]
4  [beaver, skunk, mole, hamster, squirrel, rabbi...
5  [killer+whale, blue+whale, humpback+whale, sea...
6          [bat]
7          [tiger, leopard, fox, wolf, bobcat, lion]
8  [moose, ox, sheep, buffalo, giant+panda, pig, ...
9  [dalmatian, persian+cat, german+shepherd, siam...
Name: class, dtype: object

```

To me, it looks like the clusters make pretty good sense. The large aquatic/land animals are grouped together, the flying animal is alone, the bears are together, and the household pets are grouped together.

Problem 4:

```

In [17]: from scipy.cluster.hierarchy import dendrogram, linkage
import matplotlib.pyplot as plt
%matplotlib inline
from pylab import rcParams
rcParams['figure.figsize'] = 10,10


```

```
In [18]: HC = linkage(animals_data, 'ward')
```

```

In [20]: plt.title('Hierarchical Clustering of Animals')
plt.xlabel('sample index')
plt.ylabel('distance')
dendrogram(HC, labels= classes, orientation='right')
plt.show()

```



The hierarchial clusters make sense to me as the larger land animals are grouped together, the smaller land animals are grouped together, and the aquatic animals are grouped together. In these examples, however, I believe the K-means was very comparable, especially when I see the Bat and Monkey family so similar.

0.3 Worksheet 10 - PCA and SVD

Problem 4:

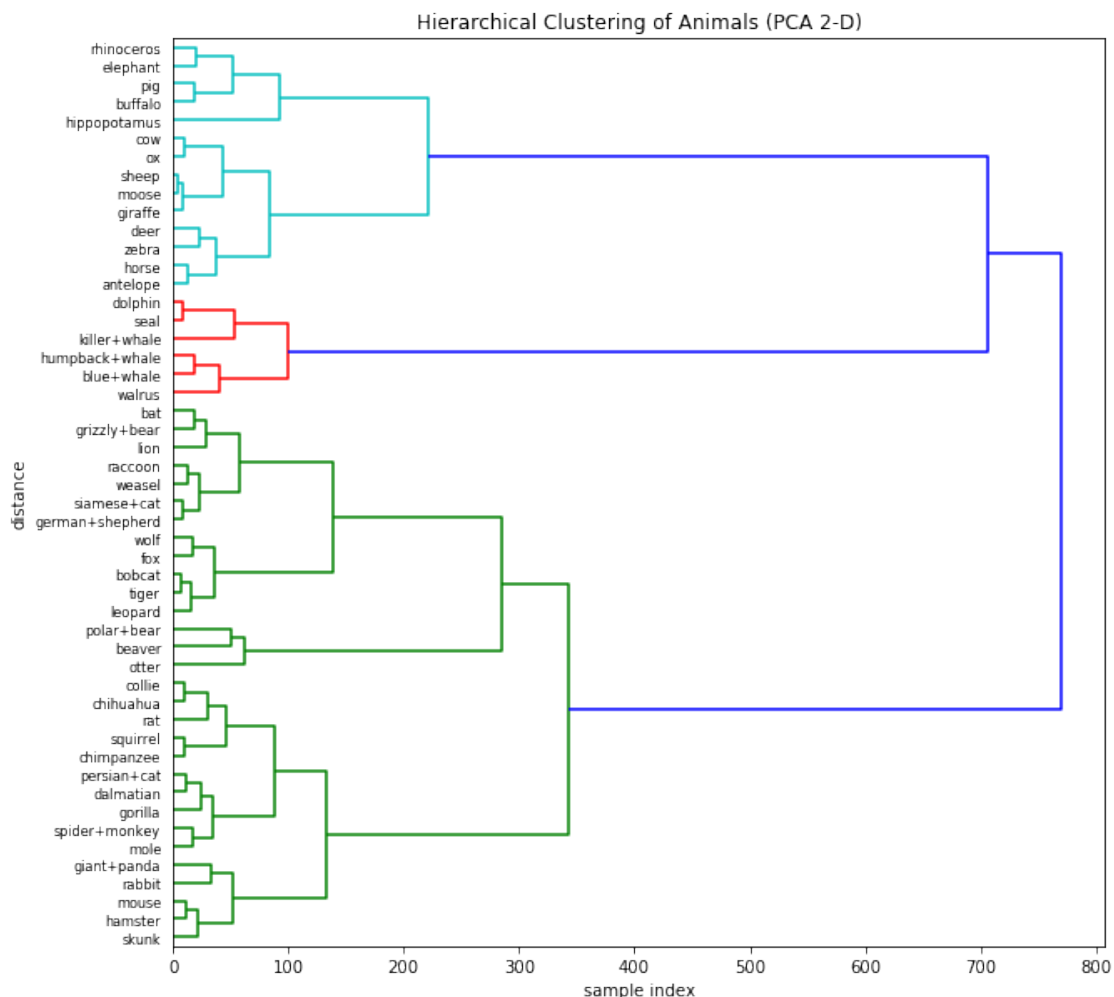
```
In [21]: from sklearn.decomposition import PCA
         pc2 = PCA(2)
         animals_data2d = pc2.fit_transform(animals_data)
         print(animals_data2d.shape)
```

(50, 2)

```
In [25]: print('Reduced to 2-D dimensionality retained',sum(pc2.explained_variance_ratio_), 'of t
```

Reduced to 2-D dimensionality retained 0.389615467763 of the original datas variance

```
In [26]: # Let us plot the hierarchial clusters again:
HC2 = linkage(animals_data2d, 'ward')
plt.title('Hierarchical Clustering of Animals (PCA 2-D)')
plt.xlabel('sample index')
plt.ylabel('distance')
dendrogram(HC2, labels= classes, orientation='right')
plt.show()
```



```
In [27]: # Plot values in 2-D
fig = plt.figure(1, figsize=(10, 10))
ax = fig.add_subplot(111)
for i, point in enumerate(animals_data2d):
    ax.annotate(classes[i], xy=point, xytext=point)
```

```
plt.scatter(animals_data2d[:,0], animals_data2d[:,1])
plt.title('PCA Projection of Animals')
```

Out[27]: <matplotlib.text.Text at 0x119db05c0>



While it does seem sensible, I think a few higher dimensions may yield more accurate results. This just seems to separate aquatic from non-aquatic animals